



SYSTEM ARCHITECTURE & IMPLEMENTATION

Multimodal AI RAG Agent

An enterprise-grade blueprint for autonomous retrieval, visual reasoning, and multi-source context fusion using modern vector databases and vision models.



Core Architectural Foundations


Unifying unstructured text, charts, tabular data, and visual imagery into a cohesive agentic decision network.

What is Multimodal Agentic RAG?

Beyond Text-Only Retrieval

Multimodal RAG combines traditional textual retrieval with visual processing and autonomous agent decision loops. It allows systems to parse complex mixed-media knowledge bases without context loss.

- **Cross-Modal Embeddings:** Connects visual features directly with text tokens in a shared vector space.
- **Vision LLM Integration:** Utilizes models like GPT-4V or Gemini Vision to interpret complex diagrams.
- **Grounding & Attribution:** Pinpoints exact document images and sub-regions for zero hallucination.



Core Supported Modalities



Text & Documents

Processes unstructured PDFs, technical specifications, legal contracts, and markdown manuals with semantic vector search.



Visuals & Diagrams

Parses flowcharts, architectural schematics, financial plots, and infographics via Vision Transformers (ViT) and OCR.



Structured Tables


Converts semi-structured HTML/CSV tables into serialized Markdown or JSON representations for exact data queries.

Multimodal Data Ingestion

Structured Ingestion Pipeline

Raw documents enter the ingestion pipeline where document layout analysis (e.g., Unstructured, Marker) separates narrative text from embedded visual objects.

Visual elements undergo automatic metadata generation via captioning models (e.g., BLIP-2) or ColPali patch-level embeddings, ensuring instant similarity search readiness in the vector store.



Autonomous Agent Decision Engine


Dynamic Route Selection

The agent evaluates user query intent and decides whether to invoke vector search, structured SQL query generators, or vision model inspection tools based on required details.

Self-Correction & Refinement


If initial retrieval returns low similarity scores or conflicting information, the agent autonomously executes query re-writing and fetches additional document patches.

Vector Database Performance



Benchmark analysis of multimodal vector embedding indexing throughput across 1M+ hybrid image-text vector embeddings.

End-to-End Query Execution



System Optimization Guardrails

-  **Patch-Level ColPali Integration:** Utilizes multi-vector representations to preserve full-page PDF layouts without losing low-level image details.
-  **Context Window Optimization:** Prioritizes visual patches with high relevance metrics to remain strictly within model token boundaries.
-  **Reciprocal Rank Fusion (RRF):** Combines sparse keyword search with dense multimodal vector retrieval for superior precision.
-  **Hallucination Guardrails:** Verifies every generated claim against retrieved source metadata prior to presenting the response to the user.

Modality Encoder Comparison

Modality	Encoder Model	Vector Dimensions	Primary Benefit
Text Documents	BGE-M3 / OpenAI v3	1536 dim	High semantic precision across long technical passages.
Images & Diagrams	CLIP / ColPali (ViT)	1024 dim (Multi-vector)	Direct alignment of visual features with language tokens.
Audio / Speech	Whisper Large v3	512 dim (Transcript)	Accurate acoustic transcription prior to vector indexing.
Tabular Data	Table Transformer (TATR)	768 dim	Preserves row-column structural relations during search.

Enterprise Application Domains



Medical Diagnostics

Correlates clinical patient notes with X-ray/MRI imagery and lab reports for decision assistance.



Industrial Engineering

Retrieves specific maintenance manual diagrams and troubleshooting steps during live field repair.



Financial Analysis


Parses annual financial disclosures containing complex embedded charts, disclosures, and balance tables.

Questions & Discussion

Multimodal AI RAG Agent System Architecture

Repository: github.com/Ahmiii-18/Custom-RAG-Agent

Image Sources



<https://mlrwd9rnffxq.i.optimole.com/cb:641c.2be21/w:1024/h:783/q:90/f:best/sm:0/https://blog.vectorize.io/wp-content/uploads/2024/10/image-12.png>

Source: [vectorize.io](https://blog.vectorize.io)



https://assets.zenml.io/webflow/64a817a2e7e2208272d1ce30/4e415291/67af0d153c9034c1e47448a9_AD_4nXe1hswTSquDZjKFP3-7u8_YYslj42-wrb5LBq73zroSs1mDDO3WaiUkQrkfcKT1il9fS1Jje2MfYITjfrSpNIKO-E2UT_agS3qJ5x_A6tdRo6omwWKdx_7V9ayMN7Zad0QrRbd5TA.png

Source: www.zenml.io